

Cloud Databases Assignment 2

Group 2
102061105 Yun-Sheng, Chang
102061212 Jian-Hao, Huang

April 5, 2017

1 Implementation

1.1 New interface Explainable

To support the *EXPLAIN* operation, we first define a new interface *Explainable*,

```
1 public interface Explainable {  
2     String getStatisticResult(int level);  
3 }
```

Explainable.java

and have every plans that would possibly be created by the *BasicQueryPlanner* implement the interface.

```
1 public class TablePlan implements Plan, Explainable  
2 public class ProductPlan implements Plan, Explainable  
3 public class SelectPlan extends ReduceRecordsPlan implements Explainable  
4 public class GroupByPlan extends ReduceRecordsPlan implements Explainable  
5 public class ProjectPlan implements Plan, Explainable  
6 public class SortPlan implements Plan, Explainable
```

Plans implementing Explainable

Three things should be done in *getStatisticResult()*:

1. indentation;
2. form the statistic result of the current plan;
3. call *p.getStatisticResult()* to obtain the statistic results of the underlying plans.

We show the implementation in *SelectPlan* as an example:

```
1 public String getStatisticResult(int level) {  
2     Explainable p;  
3     String res = new String();  
4  
5     // step 1: indentation  
6     for (int i = 0; i < level; i++)  
7         res += "\t";  
8  
9     // step 2: form the statistic result of the current plan  
10    res += "->SelectPlan pred:(" + this.pred.toString() + ") " +  
11        "(#blks=" + blocksAccessed() + ", #recs=" + recordsOutput() + ")\n";  
12  
13    // step 3: call p.getStatisticResult() to obtain the statistic result of the underlying plans  
14    p = (Explainable)this.p;  
15    res += p.getStatisticResult(level + 1);  
16  
17    return res;  
18 }
```

SelectPlan.java

The argument *level* represents the level of the current plan in the plan tree. For instance,




Figure 1: Level of plans

With the value of *level*, we will be able to do the indentation properly.

```

1 // step 1: indentation
2 for (int i = 0; i < level; i++)
3     res += "\t";

```

Indentation

Then, we can obtain the estimated number of accessed blocks and records with *blocksAccessed()* and *recordsOutput()*, respectively.

```

1 // step 2: form the statistic result of the current plan
2 res += "->SelectPlan pred:(" + this.pred.toString() + ") " +
3     "(#blks=" + blocksAccessed() + ", #recs=" + recordsOutput() + ")\n";

```

Form statistic result of the current plan

We call *p.getStatisticResult()* to get the results of the underlying plans.

```

1 // step 3: call p.getStatisticResult() to obtain the statistic result of the underlying plans
2 p = (Explainable)this.p;
3 res += p.getStatisticResult(level + 1);
4
5 return res;

```

Obtain statistic results of the underlying plans

Finally, we return the result to the upper plan.

1.2 Lexer, Parser, QueryData and BasicQueryPlanner

The keyword *explain* is added to the keyword list.

```

1 private void initKeywords() {
2     keywords = Arrays.asList("select", "from", "where", "and", "insert",
3         ...
4         "min", "max", "distinct", "group", "add", "sub", "mul", "div", "explain");
5 }

```

Lexer.java

We modify the parsing procedure to check if the query is an *EXPLAIN* query; if so, set the flag *isExplain*.

```

1 public QueryData queryCommand() {
2     boolean isExplain = false;
3     if (lex.matchKeyword("explain")) {
4         lex.eatKeyword("explain");
5         isExplain = true;
6     }
7
8     lex.eatKeyword("select");
9     ProjectList projs = projectList();
10    // parse other tokens
11
12    return new QueryData(projs.asStringSet(), tables, pred,
13        groupFields, projs.aggregationFns(), sortFields, sortDirs, isExplain);
14 }

```

Parser.java

When the query is successfully parsed, the query's information will be stored in the data structure *QueryData* and passed to the *BasicQueryPlanner*. Thus, we need to add an attribute *isExplain* to the *QueryData*.

```

1 private boolean isExplain;
2
3 public boolean isExplain() {
4     return isExplain;
5 }

```

QueryData.java

Finally, the *BasicQueryPlanner* should create the *ExplainPlan* as the top-level plan for an *EXPLAIN* query.

```

1 public Plan createPlan(QueryData data, Transaction tx) {
2     // create other plans
3
4     if (data.isExplain())
5         p = new ExplainPlan(p);
6     return p;
7 }

```

BasicQueryPlanner.java

1.3 ExplainPlan

The result of an *EXPLAIN* query contains one column, *query-plan*; we add the column to the schema in the constructor of *ExplainPlan*.

```

1 public ExplainPlan(Plan p) {
2     this.p = p;
3     this.hist = p.histogram();
4     schema.addField("query-plan", Type.VARCHAR(500));
5 }

```

ExplainPlan.java

Before the *ExplainScan* is created, we must collect the statistic result of other plans by invoking the method we previously implement: *getStatisticResult()*.

```

1 public Scan open() {
2     String str;
3
4     str = ((Explainable)p).getStatisticResult(0);
5     Scan s = p.open();
6     return new ExplainScan(s, str);
7 }

```

ExplainPlan.java

1.4 ExplainScan

ExplainScan has the following attributes:

```
1 private Scan s;  
2 private String str;  
3 private boolean first;  
4 private int numOfRec;  
5 private SpResultRecord rec;
```

ExplainScan.java

s is the scan of the underlying query.

str is the statistic results.

first is true if the client has called *beforeFirst()* and has not called *next()* yet.

numOfRec is the actual number of records of the *SELECT* query.

rec is the table storing the result of the *EXPLAIN* query.

In the constructor of *ExplainScan*, we first count the actual number of records by putting *s.next()* in a while loop and increase *numOfRec* by one every time the condition holds. Then, we append the actual number of records to the statistic result and add the final result to the table.

```
1 public ExplainScan(Scan s, String str) {  
2     s.beforeFirst();  
3     while (s.next())  
4         numOfRec++;  
5  
6     rec.setVal("query-plan",  
7         new VarcharConstant("\n" + str + "\nActual #recs: " +  
8             numOfRec, Type.VARCHAR(500)));  
9 }
```

ExplainScan.java

Since the result of the *EXPLAIN* query contains exactly one record, we set the flag *first* when *beforeFirst()* is invoked and reset it when *next()* is invoked.

```
1 public void beforeFirst() {  
2     s.beforeFirst();  
3     first = true;  
4 }  
5  
6 public boolean next() {  
7     if (first == true) {  
8         first = false;  
9         return true;  
10    }  
11    return false;  
12 }
```

ExplainScan.java

2 Results

2.1 Single table A query with WHERE

```
SQL> SELECT d_id FROM district WHERE d_w_id = 2

  d_id
-----
     1
     2
     3
     4
     5
     6
     7
     8
     9
    10

SQL> EXPLAIN SELECT d_id FROM district WHERE d_w_id = 2
|
query-plan
-----
->ProjectPlan  (#blks=3, #recs=10)
    ->SelectPlan pred:(d_w_id=2.0) (#blks=3, #recs=10)
        ->TablePlan on (district) (#blks=3, #recs=20)

Actual #recs: 10
```

2.2 Multiple tables query with WHERE

```
SQL> SELECT d_id FROM district, warehouse WHERE d_w_id = w_id

  d_id
-----
     1
     2
     3
     4
     5
     6
     7
     8
     9
    10
     1
     2
     3
     4
     5
     6
     7
     8
     9
    10

SQL> EXPLAIN SELECT d_id FROM district, warehouse WHERE d_w_id = w_id
|
query-plan
-----
->ProjectPlan  (#blks=43, #recs=20)
    ->SelectPlan pred:(d_w_id=w_id) (#blks=43, #recs=20)
        ->ProductPlan  (#blks=43, #recs=40)
            ->TablePlan on (district) (#blks=3, #recs=20)
            ->TablePlan on (warehouse) (#blks=2, #recs=2)

Actual #recs: 20
```

2.3 Query with ORDER BY

```
SQL> SELECT d_id FROM district, warehouse WHERE d_w_id = w_id ORDER BY d_id ASC
```

```
  d_id
-----
    1
    1
    2
    2
    3
    3
    4
    4
    5
    5
    6
    6
    7
    7
    8
    8
    9
    9
   10
   10
```

```
SQL> EXPLAIN SELECT d_id FROM district, warehouse WHERE d_w_id = w_id ORDER BY d_id ASC
```

query-plan

```
-----
->SortPlan  (#blks=1, #recs=20)
  ->ProjectPlan  (#blks=43, #recs=20)
    ->SelectPlan pred:(d_w_id=w_id) (#blks=43, #recs=20)
      ->ProductPlan  (#blks=43, #recs=40)
        ->TablePlan on (district) (#blks=3, #recs=20)
        ->TablePlan on (warehouse) (#blks=2, #recs=2)
```

Actual #recs: 20

2.4 Query with GROUP BY and at least one aggregation function

```
SQL> SELECT COUNT(d_id) FROM district, warehouse WHERE d_w_id = w_id GROUP BY w_id
```

```
count_of_d_id
-----
    10
    10
```

```
SQL> EXPLAIN SELECT COUNT(d_id) FROM district, warehouse WHERE d_w_id = w_id GROUP BY w_id
```

query-plan

```
-----
->ProjectPlan  (#blks=4, #recs=2)
  ->GroupByPlan  (#blks=4, #recs=2)
    ->SortPlan  (#blks=4, #recs=20)
      ->SelectPlan pred:(d_w_id=w_id) (#blks=43, #recs=20)
        ->ProductPlan  (#blks=43, #recs=40)
          ->TablePlan on (district) (#blks=3, #recs=20)
          ->TablePlan on (warehouse) (#blks=2, #recs=2)
```

Actual #recs: 2