

Cloud Databases Assignment 4

Group 2
102061105 Yun-Sheng, Chang
102061212 Jian-Hao, Huang

May 3, 2017

In this assignment, we implement the **conservative locking scheme** for concurrency control, and compare the performance result with the original **timeout-rollback locking scheme**.

1 Implementation

1.1 New Class: RecordKey

SerializableConcurrencyMgr use *String*, *BlockId* and *RecordId* as the **synchronized objects** to protect files, blocks and records from race condition, respectively; our implementation of *ConcurrencyMgr*, *ConservativeConcurrencyMgr*, use a particular class, *RecordKey*, to protect records from race condition.

```
1 private String tableName;  
2 private String field;  
3 private Constant value;  
4 private boolean isRead;  
5 private int hashCode;  
6  
7 public RecordKey(String tableName, String field, Constant value, boolean isRead) {  
8     this.tableName = tableName;  
9     this.field = field;  
10    this.value = value;  
11    this.isRead = isRead;  
12  
13    String s;  
14    s = tableName + field + value.toString();  
15    this.hashCode = s.hashCode();  
16 }
```

RecordKey.java

Although *RecordKey* has 5 member variables, only 3 of them, *tableName*, *field*, *value*, are used to represent a record. Member variable *isRead* is used to indicate which type of locks (shared or exclusive) to acquire, and regardless of the lock type, locks should always take effect on the same record if the *tableName*, *field* and *value* are identical. Member variable *hashCode* is initialized in constructor to avoid the recomputation of hash code.

```
1 @Override  
2 public boolean equals(Object obj) {  
3     if (obj == null || !(obj.getClass().equals(RecordKey.class)))  
4         return false;  
5     RecordKey r = (RecordKey) obj;  
6     return tableName.equals(r.tableName) && field.equals(r.field) &&  
7         value.equals(r.value);  
8 }  
9  
10 @Override  
11 public int hashCode() {  
12     return hashCode;  
13 }
```

RecordKey.java

1.2 LockTable

We modify *LockTable* based on the following two requirements:

1. locks can be tested for their availability (lockable or not) **outside** the *LockTable*;
2. tx with higher tx number should always acquire the same lock **after** the tx with lower tx number.

We first describe the reason for requirement 1 and how we modify the API. By definition, a transaction should **atomically** acquire all the locks it needs before its execution under the conservative locking scheme; previous implementation of *LockTable*, however, only exposes interfaces like *sLock()* and *xLock()* which invoke *wait()* immediately when the lock is unlockable, failing to make the process of acquiring all locks atomic. Thus, we define two synchronized wrapper functions in *LockTable* so that the concurrency manager can test the availability of locks before actually acquire the locks.

```
1 synchronized boolean sLockableSync(Object obj, long txNum) {
2     return sLockable(obj, txNum);
3 }
4
5 synchronized boolean xLockableSync(Object obj, long txNum) {
6     return xLockable(obj, txNum);
7 }
```

LockTable.java

Requirement 2 is in the specification of this assignment and we describe our implementation. We add a queue of type *TreeSet* to the inner class *Lockers*.

```
1 class Lockers {
2     // a variety of lockers ...
3     TreeSet<Long> queue;
4
5     Lockers() {
6         // initialize lockers ...
7         queue = new TreeSet<Long>();
8     }
9 }
```

LockTable.java

Everytime a transaction fails to atomically acquire the locks for its read/write sets, the transaction would enqueue itself to every *Lockers* that the read/write sets associate with; and when the transaction can finally acquire all the locks, it will dequeue itself before it actually acquires those locks (more detail in Section 1.4).

```
1 synchronized void enqueue(Object obj, long txNum) {
2     Lockers lks = prepareLockers(obj);
3     lks.queue.add(txNum);
4 }
5
6 synchronized void dequeue(Object obj, long txNum) {
7     Lockers lks = lockerMap.get(obj);
8     lks.queue.remove((Long) txNum);
9 }
```

LockTable.java

To forbid tx with higher tx number acquiring the same lock prior to tx with lower tx number, we modify the *Lockable()* interface, with one additional condition: **a lock is only available to the tx that has the minimal value of tx number among the queue**. We only modify the *sLockable()* and the *xLockable()* since the MGL is removed from the specification.

```
1 private boolean havePriority(Object obj, long txNum) {
2     Lockers lks = lockerMap.get(obj);
3     return lks == null || lks.queue.size() == 0 || txNum <= lks.queue.first();
4 }
5
6 private boolean sLockable(Object obj, long txNum) {
7     return (!xLocked(obj) || hasXLock(obj, txNum))
8         // unchanged conditions ...
9         && havePriority(obj, txNum);
10 }
```

```

11
12 private boolean xLockable(Object obj, long txNum) {
13     return (!sLocked(obj) || isTheOnlySLocker(obj, txNum))
14         // unchanged conditions ...
15         && havePriority(obj, txNum);
16 }

```

LockTable.java

1.3 MicrobenchmarkProc

The conservative locking scheme only works when the read/write sets are collected in advance. Thus, we construct a *List of RecordKey*, with each *RecordKey* being one record that will be accessed in a transaction.

```

1 private List<RecordKey> recordKeys = new ArrayList<RecordKey>();
2
3 public void prepare(Object... pars) {
4     paramHelper.prepareParameters(pars);
5     for(int i = 0; i < paramHelper.getReadCount(); i++){
6         Constant id = new IntegerConstant(paramHelper.getItemId(i));
7         RecordKey recKey = new RecordKey("item", "i_id", id, true);
8         recordKeys.add(recKey);
9     }
10    for(int i = 0; i < paramHelper.getWriteCount(); i++){
11        Constant id = new IntegerConstant(paramHelper.getItemId(i));
12        RecordKey recKey = new RecordKey("item", "i_id", id, false);
13        recordKeys.add(recKey);
14    }
15 }

```

MicrobenchmarkProc.java

Then, the *ConservativeConcurrencyMgr* of a transaction will try to acquire all the locks with the collected read/write sets before the transaction is actually being executed.

```

1 public SpResultSet execute() {
2     tx = VanillaDb.txMgr().newTransaction(Connection.TRANSACTION_SERIALIZABLE,
3         false);
4     ((ConservativeConcurrencyMgr) tx.concurrencyMgr()).acquireAllLocks(recordKeys);
5     // execute the transaction ...
6 }

```

MicrobenchmarkProc.java

1.4 ConservativeConcurrencyMgr

To acquire all the locks atomically, the *ConservativeConcurrencyMgr* would perform the following three steps in a critical section:

Step 1: Test if every lock is lockable; if there exist at least one lock that is unlockable (2 possible reasons: the lock is currently hold by another transaction or there exist a transaction with lower transaction number also waiting for the same lock), the *ConservativeConcurrencyMgr* would queue **all** the locks it requires and invoke *wait()*.

```

1 public void acquireAllLocks(List<RecordKey> recordKeys) {
2     boolean allLockable = false;
3     boolean queued = false;
4
5     synchronized(lockTbl) {
6         // Step 1
7         while (!allLockable) {
8             allLockable = true;
9             for (int i = 0; allLockable && i < recordKeys.size(); i++)
10                 if (!tryRecordKey(recordKeys.get(i)))
11                     allLockable = false;
12
13             if (!allLockable) {
14                 if (!queued) {
15                     enqueue(recordKeys);
16                     queued = true;

```

```

17         }
18         try {
19             lockTbl.wait();
20         } catch (Exception e) {
21             e.printStackTrace();
22         }
23     }
24 }
25 ...
26 }
27 }

```

ConservativeConcurrencyMgr.java

Step 2: When all the locks can certainly be acquired, the *ConservativeConcurrencyMgr* needs to dequeue the waiting locks if it has failed to acquire all the locks atomically at least one time.

```

1 public void acquireAllLocks(List<RecordKey> recordKeys) {
2     synchronized(lockTbl) {
3         ...
4         // Step 2
5         if (queued)
6             dequeue(recordKeys);
7         ...
8     }
9 }

```

ConservativeConcurrencyMgr.java

Step 3: Finally, acquire all the locks for the read/write sets.

```

1 public void acquireAllLocks(List<RecordKey> recordKeys) {
2     synchronized(lockTbl) {
3         ...
4         // Step 3
5         for (RecordKey r : recordKeys) {
6             if(r.isRead())
7                 readRecordKey(r);
8             else
9                 modifyRecordKey(r);
10        }
11        ...
12    }
13 }

```

ConservativeConcurrencyMgr.java

2 Results

2.1 Experiment Environment

Intel Core i7-4790 CPU @ 3.60 GHz
32 GB RAM
1TB HDD
Ubuntu Linux 16.04
Java Direct I/O (Jaydio)

2.2 Speed Up

We first compare the performance result using conservative locking scheme with the one using timeout-rollback locking scheme.

	Throughput (tx)	Average Latency (ms)
Timeout-Rollback	5788	45
Conservative	106720	28

The experiment is conducted with **50 RTEs**, **conflict ratio = 0.001**, **write ratio = 0.5** and **bufferpool size = 100000**. The speed up is roughly **18x**.

2.3 Vary Bufferpool Size

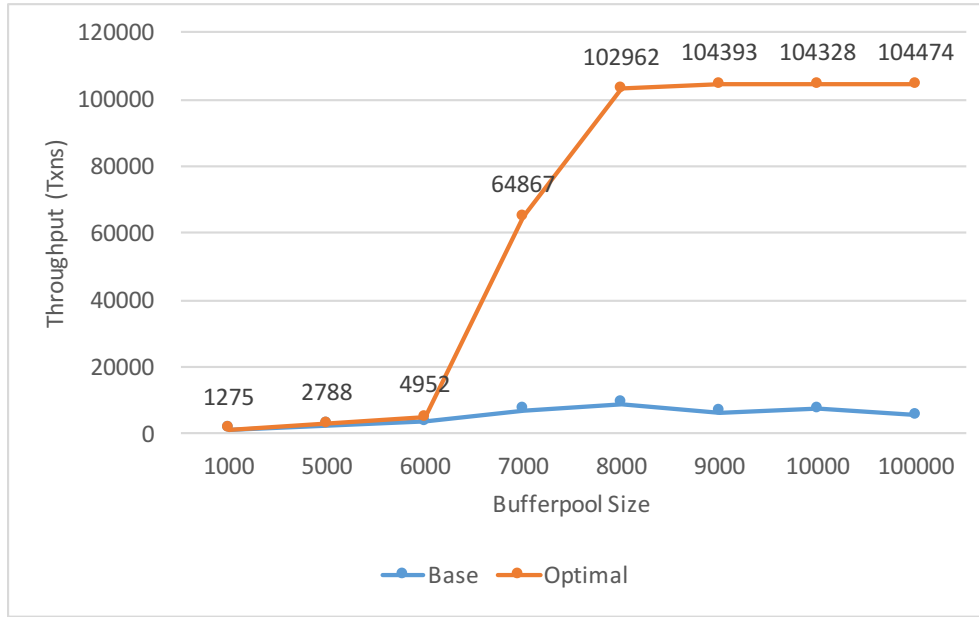


Figure 1: Throughput vs. Bufferpool Size

We vary the bufferpool size from 1000 to 100000 and we find the same trend as in assignment 3 for our conservative locking scheme: **before the match point, where the working sets just fit into the bufferpool, the throughput is low; after the match point, the throughput is high.** However, we do not find the same trend with the timeout-rollback locking scheme, we think it's because **the bottleneck is not in the size of bufferpool but in the deadlock issues.**

2.4 Vary Number of RTEs




Figure 2: Throughput vs. Number of RTEs

We vary the number of RTEs from 5 to 200 and we find that the throughput starts to drop when the number of RTEs exceeds 100. We think it's because **the more RTEs we have, the more conflict operations there will be, and naturally, the more time we will have to spend on the locks.**